

Recommender Systems Matrix Factorization


Rodrygo Santos
rodrygo@dcc.ufmg.br

Quick recap

- How will I like item x ?
 - User-based
 - What have similar users found of item x ?
 - Item-based
 - What have I found of items similar to x ?
- Key difference: neighborhoods
 - User-based: unstable, hard to precompute
 - Item-based: stable, easy to precompute


How robust is this?

- What if we have sparse data?
 - E.g., new user a , new item p
 - No available rating r_{ap}



How robust is this?

- What if we have perturbations?
 - Positively rate your own items
 - Negatively rate your competitors'



Memorize or model?

- Memory-based
 - Online “learning”
 - Online prediction
 - Lazy, instance-based learning (k -NN like)

Problems?

- Costly recommendation
 - $O(mn)$ worst case
- Poor robustness
 - Sparsity
 - Perturbations
- Overfit representation
 - I like Star Wars, you like Star Trek
 - Are we neighbors?

Memorize or model?

- Memory-based
 - Online “learning”
 - Online prediction
 - Lazy, instance-based learning (k -NN like)
- Model-based
 - Offline learning
 - Online prediction

Can we compactly represent users' tastes and item descriptions?

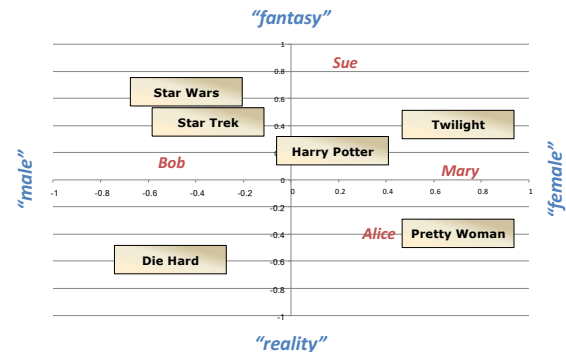
Dimensionality reduction

UF **m** G
COMPUTER
SCIENCE

- Given an $m \times n$ utility matrix
 - User-based CF is n -dimensional
 - Item-based CF is m -dimensional
 - m and n could be in the hundreds of millions
- Can we **reduce the dimensions** of the utility matrix while effectively **retaining information** about each user's preferences?

Lower dimensional projection

UF **m** G
COMPUTER
SCIENCE



Where did this come from?

UF **m** G
COMPUTER
SCIENCE

- Vocabulary mismatch problem
 - Queries are vectors over keywords
 - Documents are vectors over keywords
 - We want to match **concepts**!
- Latent semantic indexing (LSI) (Deerwester et al., ASIS 1988)
 - Represent queries and documents in a compact and robust space of latent concepts
- Actually a bit older... (like a century!)


Singular value decomposition

UF **m** G
COMPUTER
SCIENCE

- Reduce dimensionality of the problem
 - Results in small, fast model
 - Richer, denser neighbor network
- One of various matrix factorization techniques
 - Techniques based on linear systems (e.g., LU)
 - Techniques based on eigenvalues (e.g., SVD)


Factorization with SVD

UF **m** G
COMPUTER
SCIENCE



Factorization with SVD


UF **m** G
COMPUTER
SCIENCE



- "Truncated" SVD: keep k "most important" factors
 - Best rank- k approximation (by Frobenius norm)

Predictions with SVD

UF **m** G
COMPUTER
SCIENCE



$$\hat{r}_{ap} = \bar{r}_a + \frac{\sum_{b \in N} \text{sim}(a, b) \times (r_{bp} - \bar{r}_b)}{\sum_{b \in N} \text{sim}(a, b)} \approx \hat{r}_{ap} = \bar{r}_a + U_a \times \Sigma \times V_p^T$$

SVD pros

UF **m** G
COMPUTER
SCIENCE

- Prediction quality generally increases...
 - Noisy ratings filtered out
 - Nontrivial correlations detected
- ... although it may also decrease
 - Original ratings are not taken into account
- Depends on the amount of reduction
 - Normally 20 to 100 factors (Koren, KDD 2009)
 - But *it really depends* on the target domain

SVD cons

UF **m** G
COMPUTER
SCIENCE

- Lack of transparency
 - Optimal dimensions do not correspond to user-comprehensible concepts
- Missing values
 - SVD is undefined for incomplete matrices
 - Utility matrix has lots of missing values
- Computation complexity
 - SVD computation is $O(m^2n + n^3)$

Missing values


UF **m** G
COMPUTER
SCIENCE

- SVD assumes a complete matrix
 - If it's complete, we don't need a recommender!
- What to do with missing values?
 - Impute (assume they are a mean)
 - Normalize (assume they are 0)
 - Ignore!

Computational complexity

UF **m** G
COMPUTER
SCIENCE


- Standard SVD is (very) slow
 - All we need is a good approximation
- Model it as an optimization problem!



Computational complexity

UF **m** G
COMPUTER
SCIENCE

- Standard SVD is (very) slow
 - All we need is a good approximation
- Model it as an optimization problem!



Quantifying the error



- Approximated SVD
 - $\hat{r}_{ap} = U_a V_p^T$ (Σ embedded in U and V)
- Approximation error
 - $e_{ap} = r_{ap} - \hat{r}_{ap}$

Defining the problem




- We want to minimize the error over all **known ratings**, hence **ignoring missing** values
 - $\min_{U,V} \sum_{(a,p)} (r_{ap} - U_a V_p^T)^2 \quad (r_{ap} \text{ is known})$
- Problem: overfitting
 - We are learning from few examples after all
- Solution: regularization
 - $\min_{U,V} \sum_{(a,p)} (r_{ap} - U_a V_p^T)^2 + \lambda(|U_a|^2 + |V_p|^2)$

The effect of regularization



- Occam's razor (aka KISS)
 - Allow rich model where there are sufficient data
 - Shrink aggressively where data are scarce



A funky solution



Simon Funk's SVD



- One of the most interesting findings during the **Netflix Prize** came out of a blog post
 - “Ok, so here's where I tell all about how I got to be tied for third place on the Netflix prize.”
 - Simon Funk, 2006
- Incremental, iterative, and approximate way to compute the SVD using **gradient descent**

Gradient descent




- Problem
 - $\min F(\mathbf{x})$
- Solution
 - Gradient points in the direction of the greatest decrease (or increase) of a function
 - Why not step iteratively in that direction?
 - $\mathbf{x}_{i+1} = \mathbf{x}_i - \gamma_i \nabla F(\mathbf{x}_i)$
 - γ_i is the learning rate
 - Higher = faster training, lower accuracy

Gradient descent

UF **m**G
COMPUTER
SCIENCE

$$\mathbf{x}_{i+1} = \mathbf{x}_i - \gamma \nabla F(\mathbf{x}_i)$$



A simplified example

UF **m**G
COMPUTER
SCIENCE

- Goal is to decompose the left-side matrix

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} u_{11} & u_{12} \\ u_{21} & u_{22} \\ u_{31} & u_{32} \\ u_{41} & u_{42} \\ u_{51} & u_{52} \end{bmatrix} \times \begin{bmatrix} v_{11} & v_{12} & v_{13} & v_{14} & v_{15} \\ v_{21} & v_{22} & v_{23} & v_{24} & v_{25} \end{bmatrix}$$

A simplified example

UF **m**G
COMPUTER
SCIENCE

- An initial (random) guess

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix} \times \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

A simplified example

UF **m**G
COMPUTER
SCIENCE

- Not that impressive...

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \end{bmatrix}$$

- RMSE = 1.806

A simplified example

UF **m**G
COMPUTER
SCIENCE

- An improved guess

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} x & 1 \\ 1 & 1 \\ 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix} \times \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

A simplified example

UF **m**G
COMPUTER
SCIENCE

- An improved guess

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} x+1 & x+1 & x+1 & x+1 & x+1 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \end{bmatrix}$$

- Contribution of the first row $(5-(x+1))^2 + (2-(x+1))^2 + (4-(x+1))^2 + (4-(x+1))^2 + (3-(x+1))^2$

A simplified example

UF **m**G
COMPUTER
SCIENCE

- An improved guess

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} x+1 & x+1 & x+1 & x+1 & x+1 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \end{bmatrix}$$

- How to choose x to minimize the error?
– Take the gradient!

A simplified example

UF **m**G
COMPUTER
SCIENCE

- The error contribution

$$(5-(x+1))^2 + (2-(x+1))^2 + (4-(x+1))^2 + (4-(x+1))^2 + (3-(x+1))^2$$

- Simplifies to

$$(4-x)^2 + (1-x)^2 + (3-x)^2 + (3-x)^2 + (2-x)^2$$

- Taking the derivative and equating to 0

$$-2 \times ((4-x) + (1-x) + (3-x) + (3-x) + (2-x)) = 0$$

$$\therefore x = 2.6$$

A simplified example

UF **m**G
COMPUTER
SCIENCE

- Replacing $x = 2.6$

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} x+1 & x+1 & x+1 & x+1 & x+1 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \end{bmatrix}$$

A simplified example

UF **m**G
COMPUTER
SCIENCE

- Replacing $x = 2.6$

$$\begin{bmatrix} 5 & 2 & 4 & 4 & 3 \\ 3 & 1 & 2 & 4 & 1 \\ 2 & & 3 & 1 & 4 \\ 2 & 5 & 4 & 3 & 5 \\ 4 & 4 & 5 & 4 & \end{bmatrix} = \begin{bmatrix} 3.6 & 3.6 & 3.6 & 3.6 & 3.6 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \\ 2 & 2 & 2 & 2 & 2 \end{bmatrix}$$

- RMSE **reduced** from 1.806 to 1.642

Funk's SVD

UF **m**G
COMPUTER
SCIENCE

initialize matrices U, V

for each factor f :

until f has converged:

for each rating r_{ap} :

predict \hat{r}_{ap}

update U_{af}


update V_{pf}

Do we need to use
every rating in
every iteration?

Stochastic gradient descent

UF **m**G
COMPUTER
SCIENCE

- We can sample training ratings randomly
- Improves **convergence**!
– Aka **faster learning**



Summary



- Matrix factorization is dominant
 - Superior to classical nearest neighbor techniques
 - Uses a more compact memory-efficient model
- Matrix factorization is flexible
 - Can integrate multiple forms of feedback, user and item biases, temporal dynamics, and confidence levels (Koren et al., Computer 2009)
- SVD is not the only factorization approach

Writing Assignment #3



- Write a one-page summary about:
 - [Matrix factorization techniques for recommender systems](#) (Computer 2009)
by Yehuda Koren, Robert Bell, Chris Volinsky
 - Due Mon, Apr 17 @ 23:55 via Moodle